

# Hao Yan

Fairfax, VA • +1 571 354 3431 • yanhao821@gmail.com • Google Scholar • LinkedIn

## Summary

Ph.D. candidate in Computer Science at George Mason University, advised by Dr. Ziyu Yao. My research focuses on natural language processing and large language models, with a unifying goal of trustworthy, reliable, and efficient LLMs.

**Research Interests:** LLM for Code Generation, LLM Safety and Trustworthiness, Agentic AI, Mechanistic Interpretability of LLMs, LLM Explainability, Human-AI Interaction, and Persona Simulation.

## Education

### George Mason University

*Ph.D. in Computer Science, advised by Dr. Ziyu Yao*

*M.S. in Computer Science*

**Fairfax, VA**

*Aug. 2021 – Expected Dec. 2026*

*Aug. 2018 – Dec. 2020*

### Fujian Normal University

*B.E. in Electronic and Information Engineering*

**Fujian, China**

*Aug. 2013 – Jun. 2017*

## Publications (\* denotes equal contribution)

1. **Hao Yan**, Ziyu Yao. “Interpreting and Steering for Safe and Correct Code Generation”. **EMNLP 2026**, Long Paper.
  - Mechanistic-interpretability study of how code LLMs represent safe vs. vulnerable generation; built CodeSec-Pairs, a dataset of 9,342 contrastive code pairs across 5 CWEs.
  - Proposed DuoSteer, a training-free inference-time steering method that cuts the vulnerability rate of generated code by 26.9% and improves functional correctness by 7.5%.
2. **Hao Yan**, Swapneel Suhas Vaidya, Xiaokuan Zhang, Ziyu Yao. “Guiding AI to Fix Its Own Flaws: An Empirical Study on LLM-Driven Secure Code Generation”. **ACM Transactions on Software Engineering and Methodology (TOSEM) 2026**. [paper]
  - Large-scale empirical study of preventing and repairing security vulnerabilities in LLM-generated code across proprietary and open-weight models.
  - Found that directive self-generated vulnerability hints best prevent insecure code, and that detailed, actionable feedback best improves post-hoc repair.
3. Saaketh Anand Sai Kumar\*, **Hao Yan**\*, Sai Charan Perepa, Murong Yue, Ziyu Yao. “Can LLMs Simulate Personas with Reversed Performance? A Systematic Investigation for Counterfactual Instruction Following in Math Reasoning Context”. **COLM 2025 Workshop** on Social Simulation with LLMs. [paper]
  - Introduced the first benchmark for faithful counterfactual persona simulation with the Degree of Contrast (DoC) metric; showed that state-of-the-art models (e.g., o1) fail to faithfully simulate low-performing personas.
4. **Hao Yan**, Thomas D. LaToza, Ziyu Yao. “IntelliExplain: Enhancing Conversational Code Generation for Non-Professional Programmers”. **arXiv preprint**, 2024. [paper]
  - Designed a new interaction paradigm with clearer explanations and a structured feedback loop that improves non-experts’ code comprehension and task efficiency in user studies.
5. **Hao Yan**, Saurabh Srivastava, Yintao Tai, Sida I. Wang, Wen-tau Yih, Ziyu Yao. “Learning to Simulate Natural Language Feedback for Interactive Semantic Parsing”. **ACL 2023**, Long Paper. [paper]
  - Introduced the task of learning to simulate natural-language user feedback for training interactive semantic parsers; matched full-human-feedback performance on text-to-SQL even in low-data settings.
6. Saurabh Srivastava, Janit Bidhan, **Hao Yan**, et al. “Reasoning Under Constraint: How Batch Prompting Suppresses Overthinking in Reasoning Models”. **ICLR 2026 Workshop** on Logical Reasoning of Large Language Models. [paper]
  - Showed that batch prompting suppresses “overthinking” in reasoning LLMs and cuts unnecessary reasoning steps for lower-cost inference without loss of answer quality.

## Experience

### Applied Scientist Intern, Amazon (Seattle, WA)

**May 2026 – Aug. 2026**

- Researched evolutionary strategies (ES) for backpropagation-free LLM post-training at production scale.
- Designed an improved evolutionary strategy with higher task performance and better training efficiency.
- Evaluated ES post-trained models against GRPO-trained baselines and analyzed performance-efficiency trade-offs.

## **Applied Scientist Intern, Amazon** (*Seattle, WA*)

**May 2025 – Aug. 2025**

- Built a reward model that incorporates personalized user preferences into its evaluation of smart-home actions (API calls) and serves as both an automatic evaluator and the reward signal for GRPO-based post-training.
- Designed a distillation-based training strategy for the reward model and implemented the end-to-end data-processing and model-training pipeline.
- Evaluated state-of-the-art reward-modeling methods and led error analysis to identify future improvements.

## **Machine Learning Engineer Intern, Salesteq** (*Remote*)

**Jan. 2025 – Apr. 2025**

- Participated in the design of an agentic LLM framework for car sales and service workflows, using AutoGen.
- Built conversation memory and RAG-style car-inventory retrieval to ground recommendations in dealership data.
- Integrated external tools (email, scheduling APIs) via structured prompting and adaptive control logic.

## **Graduate Research & Teaching Assistant, George Mason University** (*Fairfax, VA*)

**Aug. 2021 – Present**

- Conduct research on LLM for Code Generation, LLM Safety and Trustworthiness, Agentic AI, Mechanistic Interpretability of LLMs, LLM Explainability, Human–AI Interaction, Persona Simulation, and Efficient LLM Reasoning.
- Authored multiple first-author publications on these topics.

## **Front-end Developer, Bright Sky Tech Inc.** (*Annandale, VA*)

**Jan. 2021 – Aug. 2021**

- Built and maintained a web-based online grocery platform using Angular, Firebase, TypeScript, and Ionic.

## **Professional Service**

---

### **Program Committees & Reviewing**

- **Program Committee Member:** IJCAI 2026, 2025; ECAI 2025.
- **Reviewer:** NeurIPS 2026; ICLR 2026, 2025; ACL ARR 2026, 2025; IEEE CI-NLPSoMe 2025; WLLFM 2024.

### **Invited Talks**

- “LLM Agents: Build Your Own Personal Assistant,” EPIC SPICE Intern Program, George Mason University, 2024.

### **Volunteer**

- Student Volunteer, 61st Annual Meeting of the Association for Computational Linguistics (ACL), 2023.

## **Honors & Awards**

---

- **CAHMP Summer GRA Award**, George Mason University (May 2024)
- **Graduate Student Travel Fund (GSTF)**, George Mason University (July 2023)
- **Distinguished GTA Award**, George Mason University (May 2022)
- **Graduate Students Poster Award (Honorable Mention)**, George Mason University (May 2022)
- **Outstanding GTA Award**, George Mason University (May 2020, May 2021)

## **Technical Skills**

---

**Programming:** Python, SQL, Java, JavaScript, TypeScript.

**ML & LLM Frameworks:** PyTorch, Hugging Face (Transformers, Datasets, PEFT), vLLM, AutoGen, scikit-learn, NumPy, Pandas.

**LLM Post-Training:** Supervised Fine-Tuning & LoRA, RLHF, GRPO, Evolutionary Strategies, Reward Modeling, Knowledge Distillation.

**LLM Systems & Agents:** Agentic Frameworks, Tool Use, Retrieval-Augmented Generation, Prompt Engineering.

**Evaluation & Research Methods:** Benchmark & Dataset Construction, Model Evaluation, Error Analysis, Human & LLM-Based User Studies, Model Interpretability & Explainability.

**Infrastructure:** Linux, Git, AWS, Docker, GPU/Cluster Training, Weights & Biases.